

### 6.3.4 Channel preemption

Channel preemption is enabled on a per-channel basis by setting DCHPRIn[ECP]. Channel preemption enables the executing channel's data transfers to temporarily be suspended in favor of starting a higher priority channel. After the preempting channel has completed its minor loop data transfers, the preempted channel is restored and resumes execution. After the restored channel completes one read/write sequence, it is again eligible for preemption. If any higher priority channel is requesting service, the restored channel is suspended and the higher priority channel is serviced. Nested preemption, that is, attempting to preempt a preempting channel, is not supported. After a preempting channel begins execution, it cannot be preempted. Preemption is available only when fixed arbitration is selected.

A channel's ability to preempt another channel can be disabled by setting DCHPRIn[DPA]. When a channel's preemption ability is disabled, that channel cannot suspend a lower priority channel's data transfer, regardless of the lower priority channel's ECP setting. This enables a pool of low priority, large data-moving channels to be defined. These low priority channels can be configured to not preempt each other, thus preventing a low priority channel from consuming the preempt slot normally available to a true, high priority channel.

### 6.3.5 Clocks

The following table describes the clock sources for eDMA. Please see Clock Controller Module (CCM) for clock setting, configuration and gating information.

**Table 6-5. eDMA clocks**

| Clock name | Description      |
|------------|------------------|
| edma_hclk  | Module clock     |
| ipg_clk    | Peripheral clock |

## 6.4 Initialization/application information

The following sections discuss initialization of the eDMA and programming considerations.

## 6.4.1 eDMA initialization

To initialize the eDMA:

1. Write to the CR if a configuration other than the default is desired.
2. Write the channel priority levels to the DCHPRI $n$  registers if a configuration other than the default is desired.
3. Enable error interrupts in the EEI register if desired.
4. Write the 32-byte TCD for each channel that may request service.
5. Enable any hardware service requests via the ERQ register.
6. Request channel service via either:
  - Software: setting TCD $n$ \_CSR[START]
  - Hardware: slave device asserting its eDMA peripheral request signal

After any channel requests service, a channel is selected for execution based on the arbitration and priority levels in the programming model. The eDMA engine reads the entire TCD, including the TCD control and status fields, as shown in [Table 6-6](#), for the selected channel into its internal address path module.

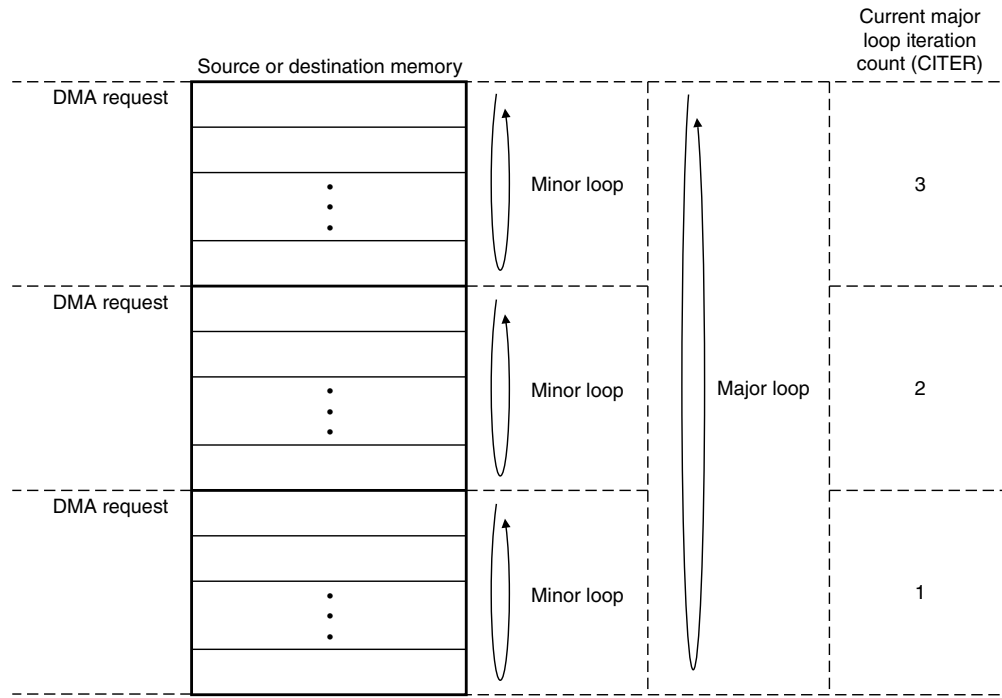
As the TCD is read, the first transfer is initiated on the system bus, unless a configuration error is detected. Transfers from the source, as defined by TCD $n$ \_SADDR, to the destination, as defined by TCD $n$ \_DADDR, continue until the number of bytes specified by TCD $n$ \_NBYTES have been transferred.

When the transfer is complete, the eDMA engine's local TCD $n$ \_SADDR, TCD $n$ \_DADDR, and TCD $n$ \_CITER are written back to the main TCD memory and any minor loop channel linking is performed, if enabled. If the major loop is exhausted, further post-processing executes, such as interrupts, major loop channel linking, and scatter/gather operations, if enabled.

**Table 6-6. TCD control and status fields**

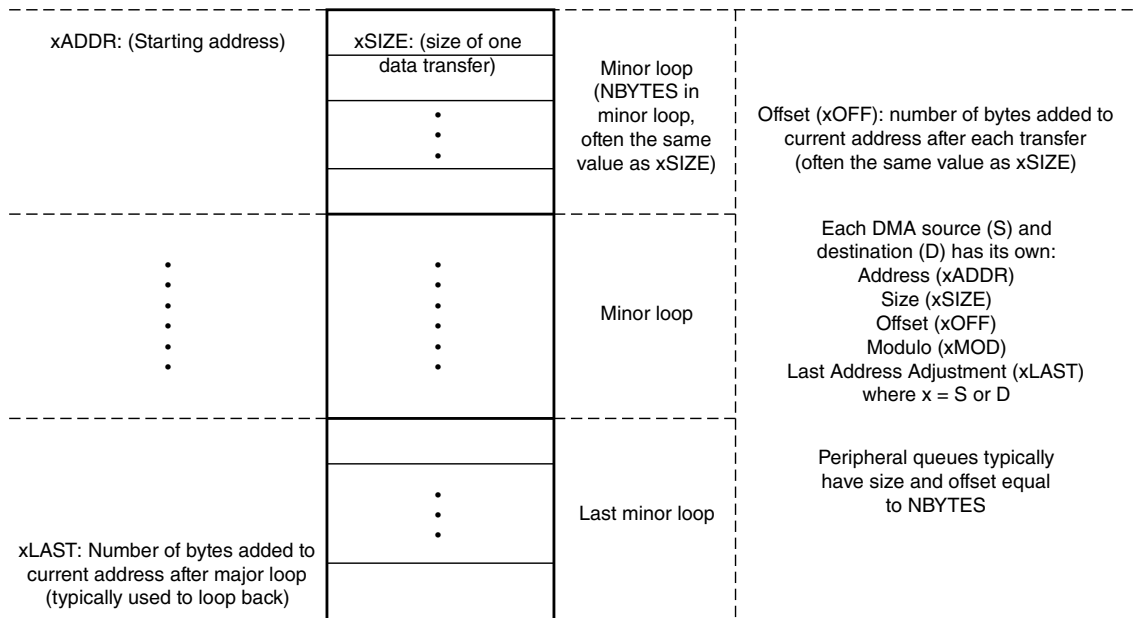
| TCD $n$ _CSR field name | Description   |
|-------------------------|---|
| START                   | Control bit to start channel explicitly when using a software-initiated DMA service (automatically cleared by hardware) |
| ACTIVE                  | Status bit indicating the channel is currently in execution   |
| DONE                    | Status bit indicating major loop completion (cleared by software when using a software-initiated DMA service)           |
| DREQ                    | Control bit to disable DMA request at end of major loop completion when using a hardware-initiated DMA service          |
| BWC                     | Control bits for throttling bandwidth control of a channel  |
| ESG                     | Control bit to enable scatter/gather feature  |
| INTHALF                 | Control bit to enable interrupt when major loop is half complete  |
| INTMAJOR                | Control bit to enable interrupt when major loop completes   |

The following figure shows how each DMA request initiates one minor-loop transfer, or iteration, without CPU intervention. DMA arbitration can occur after each minor loop, and one level of minor loop DMA preemption is allowed. The number of minor loops in a major loop is specified by the beginning iteration count (BITER).



**Figure 6-5. Example of multiple loop iterations**

The following figure lists the memory array terms and how the TCD settings interrelate.



**Figure 6-6. Memory array terms**

## 6.4.2 Programming errors

The eDMA performs various tests on the transfer control descriptor to verify consistency in the descriptor data. Most programming errors are reported on a per-channel basis with the exception of channel priority error (ES[CPE]).

For all error types other than group or channel priority errors, the channel number causing the error is recorded in the Error Status register (DMAx\_ES). If the error source is not removed before the next activation of the problem channel, the error is detected and recorded again.

Channel priority errors are identified within a group after that group has been selected as the active group. For example:

1. The eDMA is configured for fixed group and fixed channel arbitration modes.
2. Group 1 is the highest priority and all channels are unique in that group.
3. Group 0 is the next highest priority and has two channels with the same priority level.
4. If Group 1 has any service requests, those requests will be executed.
5. After all of Group 1 requests have completed, Group 0 will be the next active group.
6. If Group 0 has a service request, then an undefined channel in Group 0 will be selected and a channel priority error will occur.
7. This repeats until the all of Group 0 requests have been removed or a higher priority Group 1 request comes in.

In this sequence, for item 2, the eDMA acknowledge lines will assert only if the selected channel is requesting service via the eDMA peripheral request signal. If interrupts are enabled for all channels, the user will get an error interrupt, but the channel number for the ERR register and the error interrupt request line may be wrong because they reflect the selected channel. A group priority error is global and any request in any group will cause a group priority error.

If priority levels are not unique, when any channel requests service, a channel priority error is reported. The highest channel/group priority with an active request is selected, but the lowest numbered channel with that priority is selected by arbitration and executed by the eDMA engine. The hardware service request handshake signals, error interrupts, and error reporting are associated with the selected channel.

## 6.4.3 Arbitration mode considerations

This section discusses arbitration considerations for the eDMA.

### 6.4.3.1 Fixed group arbitration, Fixed channel arbitration

In this mode, the channel service request from the highest priority channel in the highest priority group is selected to execute. If the eDMA is programmed so that the channels within one group use "fixed" priorities, and that group is assigned the highest "fixed" priority of all groups, that group can take all the bandwidth of the eDMA controller. That is, no other groups will be serviced if there is always at least one DMA request pending on a channel in the highest priority group when the controller arbitrates the next DMA request. The advantage of this scenario is that latency can be small for channels that need to be serviced quickly. Preemption is available in this scenario only.

### 6.4.3.2 Fixed group arbitration, Round-robin channel arbitration

The highest priority group with a request will be serviced. Lower priority groups will be serviced if no pending requests exist in the higher priority groups.

Within each group, channels are serviced starting with the highest channel number and rotating through to the lowest channel number without regard to the channel priority levels assigned within the group.

This scenario could cause the same bandwidth consumption problem as indicated in [Fixed group arbitration, Fixed channel arbitration](#), but all the channels in the highest priority group will be serviced. Service latency will be short on the highest priority group, but could potentially be very much longer as the group priority decreases.

## 6.4.4 DMA transfer examples

This section presents examples of how to perform DMA transfers with the eDMA.

### 6.4.4.1 Single request

To perform a simple transfer of  $n$  bytes of data with one activation, set the major loop to one ( $\text{TCDn\_CITER} = \text{TCDn\_BITER} = 1$ ). The data transfer begins after the channel service request is acknowledged and the channel is selected to execute. After the transfer is complete,  $\text{TCDn\_CSR}[\text{DONE}]$  is set and an interrupt generates if properly enabled.

For example, the following TCD entry is configured to transfer 16 bytes of data. The eDMA is programmed for one iteration of the major loop transferring 16 bytes per iteration. The source memory has an 8-bit memory port located at 0x1000. The

destination memory has a 32-bit port located at 0x2000. The address offsets are programmed in increments to match the transfer size: one byte for the source and four bytes for the destination. The final source and destination addresses are adjusted to return to their beginning values.

```
TCDn_CITER = TCDn_BITER = 1
TCDn_NBYTES = 16
TCDn_SADDR = 0x1000
TCDn_SOFF = 1
TCDn_ATTR[SSIZE] = 0
TCDn_SLAST = -16
TCDn_DADDR = 0x2000
TCDn_DOFF = 4
TCDn_ATTR[DSIZE] = 2
TCDn_DLAST_SGA = -16
TCDn_CSR[INTMAJOR] = 1
TCDn_CSR[START] = 1 (Should be written last after all other fields have been initialized)
All other TCDn fields = 0
```

This generates the following event sequence:

1. User write to the TCDn\_CSR[START] bit requests channel service.
2. The channel is selected by arbitration for servicing.
3. eDMA engine writes: TCDn\_CSR[DONE] = 0, TCDn\_CSR[START] = 0, TCDn\_CSR[ACTIVE] = 1.
4. eDMA engine reads: channel TCD data from local memory to internal register file.
5. The source-to-destination transfers are executed as follows:
  - a. Read byte from location 0x1000, read byte from location 0x1001, read byte from 0x1002, read byte from 0x1003.
  - b. Write 32 bits to location 0x2000 → first iteration of the minor loop.
  - c. Read byte from location 0x1004, read byte from location 0x1005, read byte from 0x1006, read byte from 0x1007.
  - d. Write 32 bits to location 0x2004 → second iteration of the minor loop.
  - e. Read byte from location 0x1008, read byte from location 0x1009, read byte from 0x100A, read byte from 0x100B.
  - f. Write 32 bits to location 0x2008 → third iteration of the minor loop.
  - g. Read byte from location 0x100C, read byte from location 0x100D, read byte from 0x100E, read byte from 0x100F.
  - h. Write 32 bits to location 0x200C → last iteration of the minor loop → major loop complete.
6. The eDMA engine writes: TCDn\_SADDR = 0x1000, TCDn\_DADDR = 0x2000, TCDn\_CITER = 1 (TCDn\_BITER).
7. The eDMA engine writes: TCDn\_CSR[ACTIVE] = 0, TCDn\_CSR[DONE] = 1, INT[n] = 1.
8. The channel retires and the eDMA goes idle or services the next channel.

### 6.4.4.2 Multiple requests

The following example transfers 32 bytes via two hardware requests, but is otherwise the same as the previous example. The only fields that change are the major loop iteration count and the final address offsets. The eDMA is programmed for two iterations of the major loop, transferring 16 bytes per iteration. After the channel's hardware requests are enabled in the ERQ register, the slave device initiates channel service requests.

```
TCDn_CITER = TCDn_BITER = 2
TCDn_SLAST = -32
TCDn_DLAST_SGA = -32
```

This would generate the following sequence of events:

1. First hardware, that is, the eDMA peripheral, requests channel service.
2. The channel is selected by arbitration for servicing.
3. eDMA engine writes: TCDn\_CSR[DONE] = 0, TCDn\_CSR[START] = 0, TCDn\_CSR[ACTIVE] = 1.
4. eDMA engine reads: channel TCDn data from local memory to internal register file.
5. The source to destination transfers are executed as follows:
  - a. Read byte from location 0x1000, read byte from location 0x1001, read byte from 0x1002, read byte from 0x1003.
  - b. Write 32 bits to location 0x2000 → first iteration of the minor loop.
  - c. Read byte from location 0x1004, read byte from location 0x1005, read byte from 0x1006, read byte from 0x1007.
  - d. Write 32 bits to location 0x2004 → second iteration of the minor loop.
  - e. Read byte from location 0x1008, read byte from location 0x1009, read byte from 0x100A, read byte from 0x100B.
  - f. Write 32 bits to location 0x2008 → third iteration of the minor loop.
  - g. Read byte from location 0x100C, read byte from location 0x100D, read byte from 0x100E, read byte from 0x100F.
  - h. Write 32 bits to location 0x200C → last iteration of the minor loop.
6. eDMA engine writes: TCDn\_SADDR = 0x1010, TCDn\_DADDR = 0x2010, TCDn\_CITER = 1.
7. eDMA engine writes: TCDn\_CSR[ACTIVE] = 0.
8. The channel retires → one iteration of the major loop. The eDMA goes idle or services the next channel.
9. Second hardware, that is, eDMA peripheral, requests channel service.
10. The channel is selected by arbitration for servicing.
11. eDMA engine writes: TCDn\_CSR[DONE] = 0, TCDn\_CSR[START] = 0, TCDn\_CSR[ACTIVE] = 1.
12. eDMA engine reads channel TCD data from local memory to internal register file.
13. The source to destination transfers are executed as follows:

- a. Read byte from location 0x1010, read byte from location 0x1011, read byte from 0x1012, read byte from 0x1013.
  - b. Write 32 bits to location 0x2010 → first iteration of the minor loop.
  - c. Read byte from location 0x1014, read byte from location 0x1015, read byte from 0x1016, read byte from 0x1017.
  - d. Write 32 bits to location 0x2014 → second iteration of the minor loop.
  - e. Read byte from location 0x1018, read byte from location 0x1019, read byte from 0x101A, read byte from 0x101B.
  - f. Write 32 bits to location 0x2018 → third iteration of the minor loop.
  - g. Read byte from location 0x101C, read byte from location 0x101D, read byte from 0x101E, read byte from 0x101F.
  - h. Write 32 bits to location 0x201C → last iteration of the minor loop → major loop complete.
14. eDMA engine writes:  $TCDn\_SADDR = 0x1000$ ,  $TCDn\_DADDR = 0x2000$ ,  $TCDn\_CITER = 2$  ( $TCDn\_BITER$ ).
  15. eDMA engine writes:  $TCDn\_CSR[ACTIVE] = 0$ ,  $TCDn\_CSR[DONE] = 1$ ,  $INT[n] = 1$ .
  16. The channel retires → major loop complete. The eDMA goes idle or services the next channel.

### 6.4.4.3 Using the modulo feature

The modulo feature of the eDMA provides the ability to implement a circular data queue in which the size of the queue is a power of 2. MOD is a 5-bit field for the source and destination in the TCD, and it specifies which lower address bits increment from their original value after the address+offset calculation. All upper address bits remain the same as in the original value. A setting of zero for this field disables the modulo feature.

The following table shows how the transfer addresses are specified based on the setting of the MOD field. Here a circular buffer is created where the address wraps to the original value while the 28 upper address bits (0x1234567x) retain their original value. In this example the source address is set to 0x12345670, the offset is set to 4 bytes, and the MOD field is set to 4, allowing for a  $2^4$  byte (16-byte) size queue.

**Table 6-7. Modulo example**

| Transfer number | Address    |
|-----------------|------------|
| 1               | 0x12345670 |
| 2               | 0x12345674 |
| 3               | 0x12345678 |
| 4               | 0x1234567C |

*Table continues on the next page...*



**Table 6-7. Modulo example (continued)**

| Transfer number | Address    |
|-----------------|------------|
| 5               | 0x12345670 |
| 6               | 0x12345674 |

## 6.4.5 Monitoring transfer descriptor status

This section discusses how to monitor eDMA status.

### 6.4.5.1 Testing for minor loop completion

There are two methods to test for minor loop completion when using software-initiated service requests. The first is to read the `TCDn_CITER` field and test for a change. Another method may be extracted from the sequence shown below. The second method is to test `TCDn_CSR[START]` and `TCDn_CSR[ACTIVE]`. The minor-loop-complete condition is indicated by both bits reading zero after `TCDn_CSR[START]` was set. Polling the `TCDn_CSR[ACTIVE]` bit may be inconclusive, because the active status may be missed if the channel execution is short in duration.

The TCD status bits execute the following sequence for a software activated channel:

| Stage | TCDn_CSR bits |        |      | State  |
|-------|---------------|--------|------|--|
|       | START         | ACTIVE | DONE |  |
| 1     | 1             | 0      | 0    | Channel service request via software             |
| 2     | 0             | 1      | 0    | Channel is executing                             |
| 3a    | 0             | 0      | 0    | Channel has completed the minor loop and is idle |
| 3b    | 0             | 0      | 1    | Channel has completed the major loop and is idle |

The best method to test for minor-loop completion when using service requests initiated by hardware, that is, peripherals, is to read the `TCDn_CITER` field and test for a change. The hardware request and acknowledge handshake signals are not visible in the programmer's model.

The TCD status bits execute the following sequence for a hardware-activated channel:

| Stage | TCD $n$ _CSR bits |        |      | State  |
|-------|-------------------|--------|------|--|
|       | START             | ACTIVE | DONE |  |
| 1     | 0                 | 0      | 0    | Channel service request via hardware (peripheral request asserted) |
| 2     | 0                 | 1      | 0    | Channel is executing   |
| 3a    | 0                 | 0      | 0    | Channel has completed the minor loop and is idle                   |
| 3b    | 0                 | 0      | 1    | Channel has completed the major loop and is idle                   |

For both activation types, the major-loop-complete status is explicitly indicated via the TCD $n$ \_CSR[DONE] bit.

The TCD $n$ \_CSR[START] bit is cleared automatically when the channel begins execution regardless of how the channel activates.

#### 6.4.5.2 Reading the transfer descriptors of active channels

The eDMA reads back the true TCD $n$ \_SADDR, TCD $n$ \_DADDR, and TCD $n$ \_NBYTES values if read when a channel executes. The true values of SADDR, DADDR, and NBYTES are the values the eDMA engine currently uses in its internal register file and not the values in the TCD local memory for that channel. The addresses, SADDR and DADDR, and NBYTES, which decrement to zero as the transfer progresses, can give an indication of the progress of the transfer. All other values are read back from the TCD local memory.

#### 6.4.5.3 Checking channel preemption status

Preemption is available only when fixed arbitration is selected for both group and channel arbitration modes. A preemptive situation is one in which a preempt-enabled channel runs and a higher priority request becomes active. When the eDMA engine is not operating in fixed group, fixed channel arbitration mode, determination of the actively running relative priority outstanding requests become undefined. Channel and/or group priorities are treated as equal, that is, constantly rotating, when Round-Robin Arbitration mode is selected.

The TCD $n$ \_CSR[ACTIVE] bit for the preempted channel remains asserted throughout the preemption. The preempted channel is temporarily suspended while the preempting channel executes one major loop iteration. If two TCD $n$ \_CSR[ACTIVE] bits are set simultaneously in the global TCD map, a higher priority channel is actively preempting a lower priority channel.

## 6.4.6 Channel linking

Channel linking (or chaining) is a mechanism where one channel sets the TCD<sub>n</sub>\_CSR[START] bit of another channel (or itself), thus initiating a service request for that channel. When properly enabled, the EDMA engine automatically performs this operation at the major or minor loop completion.

The minor loop channel linking occurs at the completion of the minor loop (or one iteration of the major loop). The TCD<sub>n</sub>\_CITER[ELINK] field determines whether a minor loop link is requested. When enabled, the channel link is made after each iteration of the major loop except for the last. When the major loop is exhausted, only the major loop channel link fields are used to determine if a channel link should be made. For example, the initial fields of:

```
TCDn_CITER[ELINK] = 1
TCDn_CITER[LINKCH] = 0xC
TCDn_CITER[CITER] value = 0x4
TCDn_CSR[MAJOR_ELINK] = 1
TCDn_CSR[MAJOR_LINKCH] = 0x3
```

executes as:

1. Minor loop done → set TCD2\_CSR[START] bit.
2. Minor loop done → set TCD2\_CSR[START] bit.
3. Minor loop done → set TCD2\_CSR[START] bit.
4. Minor loop done, major loop done → set TCD3\_CSR[START] bit.

When minor loop linking is enabled (TCD<sub>n</sub>\_CITER[ELINK] = 1), the TCD<sub>n</sub>\_CITER[CITER] field uses a nine-bit vector to form the current iteration count. When minor loop linking is disabled (TCD<sub>n</sub>\_CITER[ELINK] = 0), the TCD<sub>n</sub>\_CITER[CITER] field uses a 15-bit vector to form the current iteration count. The bits associated with the TCD<sub>n</sub>\_CITER[LINKCH] field are concatenated onto the CITER value to increase the range of the CITER.

### Note

The TCD<sub>n</sub>\_CITER[ELINK] bit and the TCD<sub>n</sub>\_BITER[ELINK] bit must be equal or a configuration error is reported. The CITER and BITER vector widths must be equal to calculate the major loop, half-way done interrupt point.

The following table summarizes how a DMA channel can link to another DMA channel, that is, use another channel's TCD, at the end of a loop.

**Table 6-8. Channel linking parameters**

| Desired link behavior     | TCD control field name | Description  |
|---------------------------|------------------------|--|
| Link at end of minor loop | CITER[ELINK]           | Enable channel-to-channel linking on minor loop completion (current iteration) |
|                           | CITER[LINKCH]          | Link channel number when linking at end of minor loop (current iteration)      |
| Link at end of major loop | CSR[MAJOR_ELINK]       | Enable channel-to-channel linking on major loop completion                     |
|                           | CSR[MAJOR_LINKCH]      | Link channel number when linking at end of major loop                          |

## 6.4.7 Dynamic programming

This section provides recommended methods to change the programming model during channel execution.

### 6.4.7.1 Dynamically changing the channel priority

The following two options are recommended for dynamically changing channel priority levels:

1. Switch to Round-Robin Channel Arbitration mode, change the channel priorities, then switch back to Fixed Arbitration mode
2. Disable all the channels, change the channel priorities, then enable the appropriate channels.

### 6.4.7.2 Dynamic channel linking

Dynamic channel linking is the process of setting [TCDn\\_CSR\[MAJORELINK\]](#) during channel execution (see the diagram in [TCD structure](#)). This field is read from the TCD local memory at the end of channel execution, thus enabling you to enable the feature during channel execution.

Because you can change the configuration during execution, a coherency model is needed. Consider the scenario where you attempt to execute a dynamic channel link by enabling [TCDn\\_CSR\[MAJORELINK\]](#) at the same time the eDMA engine is retiring the channel. [TCDn\\_CSR\[MAJORELINK\]](#) would be set in the programmer's model, but it would be unclear whether the actual link was made before the channel retired.

The following coherency model is recommended when executing a dynamic channel link request.

1. Write one to TCDn\_CSR[MAJORELINK].
2. Read back TCDn\_CSR[MAJORELINK].
3. Test the TCDn\_CSR[MAJORELINK] request status:
  - If TCDn\_CSR[MAJORELINK] = 1, the dynamic link attempt was successful.
  - If TCDn\_CSR[MAJORELINK] = 0, the attempted dynamic link did not succeed (the channel was already retiring).

For this request, the TCD local memory controller forces TCDn\_CSR[MAJORELINK] to zero on any writes to a channel's TCD.word7 after that channel's TCD.done bit is set, indicating the major loop is complete.

### NOTE

You must clear [TCDn\\_CSR\[DONE\]](#) before writing TCDn\_CSR[MAJORELINK]. The eDMA engine automatically clears TCDn\_CSR[DONE] after a channel begins execution.

#### 6.4.7.3 Dynamic scatter/gather

Scatter/gather is the process of automatically loading a new TCD into a channel. It enables a DMA channel to use multiple TCDs; this enables a DMA channel to scatter the DMA data to multiple destinations or gather it from multiple sources. When scatter/gather is enabled and the channel has finished its major loop, a new TCD is fetched from system memory and loaded into that channel's descriptor location in eDMA programmer's model, thus replacing the current descriptor.

Because you are able to change the configuration during execution, a coherency model is needed. Consider the scenario where you attempt to execute a dynamic scatter/gather operation by enabling the [TCDn\\_CSR\[ESG\]](#) bit at the same time the eDMA engine is retiring the channel. The ESG bit would be set in the programmer's model, but it would be unclear whether the actual scatter/gather request was honored before the channel retired.

Two methods for this coherency model are shown in the following subsections. Method 1 has the advantage of reading the MAJORLINKCH field and the ESG bit with a single read. For both dynamic channel linking and scatter/gather requests, the TCD local memory controller forces the TCD MAJOR[ELINK] and ESG bits to zero on any writes to a channel's TCD word 7 if that channel's TCD[DONE] bit is set, indicating the major loop is complete.

**NOTE**

The user must clear the [TCDn\\_CSR\[DONE\]](#) bit before writing the MAJORELINK or ESG bits. The TCDn\_CSR[DONE] bit is cleared automatically by the eDMA engine after a channel begins execution.

**6.4.7.3.1 Method 1 (channel not using major loop channel linking)**

For a channel not using major loop channel linking, the coherency model described here may be used for a dynamic scatter/gather request.

When [TCDn\\_CSR\[MAJORELINK\]](#) is zero, TCDn\_CSR[MAJORLINKCH] is not used by the eDMA. In this case, MAJORLINKCH may be used for other purposes. This method uses the MAJORLINKCH field as a TCD identification (ID).

1. When the descriptors are built, write a unique TCD ID in TCDn\_CSR[MAJORLINKCH] for each TCD associated with a channel using dynamic scatter/gather.
2. Write one to [TCDn\\_CSR\[DREQ\]](#).

Should a dynamic scatter/gather attempt fail, setting the DREQ bit prevents a future hardware activation of the channel. This stops the channel from executing with a destination address (DADDR) that was calculated using a scatter/gather address (written in the next step) instead of a DLAST\_SGA final offset value.

3. Write the [TCDn\\_DLASTSGA](#) register with the scatter/gather address.
4. Write one to TCDn\_CSR[ESG].
5. Read back the 16-bit TCD control/status field.
6. Test the ESG request status and MAJORLINKCH value in the TCDn\_CSR register:

If ESG = 1, the dynamic link attempt was successful.

If ESG = 0 and MAJORLINKCH (ID) did not change, the attempted dynamic link did not succeed (the channel was already retiring).

If ESG = 0 and MAJORLINKCH (ID) changed, the dynamic link attempt was successful (the new TCD's ESG value cleared the ESG bit).

**6.4.7.3.2 Method 2 (channel using major loop channel linking)**

For a channel using major loop channel linking, the coherency model described here may be used for a dynamic scatter/gather request. This method uses the TCD[DLAST\_SGA] field as a TCD identification (ID).

1. Write one to [TCDn\\_CSR\[DREQ\]](#).

Should a dynamic scatter/gather attempt fail, setting TCDn\_CSR[DREQ] prevents a future hardware activation of the channel. This stops the channel from executing with a destination address (DADDR) that was calculated using a scatter/gather address (written in the next step) instead of a DLAST\_SGA final offset value.

2. Write the **TCDn\_DLAST\_SGA** register with the scatter/gather address.
3. Write one to TCDn\_CSR[ESG].
4. Read back TCDn\_CSR[ESG].
5. Test the ESG request status:

If ESG = 1, the dynamic link attempt was successful.

If ESG = 0, read the 32-bit TCDn\_DLAST\_SGA field.

If ESG = 0 and TCDn\_DLAST\_SGA did not change, the attempted dynamic link did not succeed (the channel was already retiring).

If ESG = 0 and TCDn\_DLAST\_SGA changed, the dynamic link attempt was successful (the new TCD's ESG value cleared the ESG bit).

### 6.4.8 Suspend/resume a DMA channel with active hardware service requests

The DMA enables you to move data from memory or peripheral registers to another location in memory or peripheral registers without CPU interaction. After the DMA and peripherals have been configured and are active, it is rare to suspend a peripheral's service request dynamically. In this scenario, there are certain restrictions to disabling a DMA hardware service request. For coherency, a specific procedure must be followed. This section provides guidance on how to coherently suspend and resume a Direct Memory Access (DMA) channel when the DMA is triggered by a slave module such as the Serial Peripheral Interface (SPI), ADC, or other module.

#### 6.4.8.1 Suspend an active DMA channel

To suspend an active DMA channel:

1. Stop the DMA service request at the peripheral first. Confirm it has been disabled by reading back the appropriate register in the peripheral.
2. Check the DMA's Hardware Request Status register (DMA\_HRSn) to ensure there is no service request to the DMA channel being suspended. Then disable the hardware service request by clearing the ERQ bit on appropriate DMA channel.